

# Operating System Basics

## CS 111

### Operating Systems

### Peter Reiher

# Outline

- Important properties for an operating system
- Critical abstractions for operating systems
- System services

# Important OS Properties

- For real operating systems built and used by real people
- What's most important depends on who you are talking about
  - Users
  - Service providers
  - Application developers
  - OS developers
- All are important clients for operating systems

# Reliability

- Your OS really should never crash
  - Since it takes everything else down with it
- But also need dependability in a different sense
  - The OS must be depended on to behave as it's specified
  - Nobody wants surprises from their operating system
  - Since the OS controls everything, unexpected behavior could be arbitrarily bad

# Performance

- A loose goal
- The OS must perform well in critical situations
- But optimizing the performance of all OS operations not always critical
- Nothing can take too long
- But if something is “fast enough,” adding complexity to make it faster not worthwhile
  - Often overlooked by OS researchers and developers

# Upward Compatibility

- People want new releases of an OS
  - New features, bug fixes, enhancements
- People also fear new releases of an OS
  - OS changes can break old applications
- What makes the compatibility issue manageable?
  - Stable interfaces

# Stable Interfaces

- Designers should start with well specified Application Interfaces
  - Must keep them stable from release to release
- Application developers should only use committed interfaces
  - Don't use undocumented features or erroneous side effects

# APIs

- Application Program Interfaces
  - A source level interface, specifying:
    - Include files, data types, constants
    - Macros, routines and their parameters
- A basis for software portability
  - Recompile program for the desired architecture
  - Linkage edit with OS-specific libraries
  - Resulting binary runs on that architecture and OS
- An API compliant program will compile & run on any compliant system
  - APIs are primarily for programmers



# ABIs

- Application Binary Interfaces
  - A binary interface, specifying:
    - Dynamically loadable libraries (DLLs)
    - Data formats, calling sequences, linkage conventions
  - The binding of an API to a hardware architecture
- A basis for binary compatibility
  - One binary serves all customers for that hardware
    - E.g. all x86 Linux/BSD/MacOS/Solaris/...
- An ABI compliant program will run (unmodified) on any compliant system
- ABIs are primarily for users

# Maintainability

- Operating systems have very long lives
  - Solaris, the “new kid on the block,” came out in 1993
  - Even smart phone OSes have roots in the 80s or 90s
- Basic requirements will change many times
- Support costs will dwarf initial development
- This makes maintainability critical
- Aspects of maintainability:
  - Understandability
  - Modularity/modifiability
  - Testability

# Maintainability: Understandability

- Code must be learnable by mortals
  - It will not be maintained by the original developers
  - New people must be able to come up to speed
- Code must be well organized
  - Nobody can understand 1 million lines of random code
  - It must have understandable, hierarchical structure
- Documentation
  - High level structure, and organizing principles
  - Functionality, design, and rationale for modules
  - How to solve common problems

# Why a Hierarchical Structure?

- Not absolutely necessary, but . . .
- Hierarchical layers usually understandable without completely understanding the implementation
- Expansion of one sub-system in a hierarchy usually understandable w/out understanding the expansion of other sub-systems
- Other structures tend not to have those advantages

# Maintainability: Modularity and Modifiability

- Modules must be understandable in isolation
  - Modules should perform coherent functions
  - Well-specified interfaces for each module
  - Implementation details hidden within module
  - Inter-module dependencies should be few/simple/clean
- Modules must be independently changeable
  - Lots of side effects mean lots of bugs
  - Changes to one module should not affect others
- Keep It Simple Stupid
  - Costs of complexity usually outweigh the rewards

# Side Effects

- A *side effect* is a situation where an action in one object has non-obvious consequences
  - Perhaps even to other objects
  - Generally not following the interface specification
- Side effects often happen when state is shared between seemingly independent modules and functions
- Side effects lead to unexpected behaviors
- And the resulting bugs can be hard to find

# Maintainability: Testability

- OS must work, so its developers must test it
- Thorough testing is key to reliability
  - All modules must be thoroughly testable
  - Most modules should be testable in isolation
- Testability must be designed in from the start
  - Observability of internal state
  - Triggerability of all operations and situations
  - Isolability of functionality
- Testing must be automated
  - Functionality, regression, performance,
  - Stress testing, error handling handling

# Automated Testing

- Why is it important that testing be automated?
- Automated tests can be run often (e.g. after every change) with very little cost or effort
- Automatically executed tests are much more likely to be run completely and correctly every time
- And discrepancies are much more likely to be noted and reported



# Cost of Development

- Another area where simplicity wins
- If it's simple, it will be quicker and cheaper to build
- Even better, there will be fewer bugs
  - And thus less cost for bug fixes
- And changing/extending it will be cheaper
- Low cost development usually implies speedy development
  - Quicker time to market

# Critical OS Abstractions

- One of the main roles of an operating system is to provide abstract services
  - Services that are easier for programs and users to work with
- What are the important abstractions an OS provides?

# Abstractions of Memory

- Many resources used by programs and people relate to data storage
  - Variables
  - Chunks of allocated memory
  - Files
  - Database records
  - Messages to be sent and received
- These all have some similar properties

# The Basic Memory Operations

- Regardless of level or type, memory abstractions support a couple of operations
  - WRITE(name, value)
    - Put a value into a memory location specified by name
  - value <- READ(name)
    - Get a value out of a memory location specified by name
- Seems pretty simple
- But going from a nice abstraction to a physical implementation can be complex

# Some Complicating Factors

- Persistent vs. transient memory
- Size of operations
  - Size the user/application wants to work with
  - Size the physical device actually works with
- Coherence and atomicity
- Latency
- Same abstraction might be implemented with many different physical devices
  - Possibly of very different types

# Where Do the Complications Come From?

- At the bottom, the OS doesn't have abstract devices with arbitrary properties
- It has particular physical devices
  - With unchangeable, often inconvenient, properties
- The core OS abstraction problem:
  - Creating the abstract device with the desirable properties from the physical device without them

# An Example

- A typical file
- We can read or write the file
- We can read or write arbitrary amounts of data
- If we write the file, we expect our next read to reflect the results of the write
  - *Coherence*
- If there are several reads/writes to the file, we expect each to occur in some order
  - With respect to the others

# What Is Implementing the File?

- Most commonly a hard disk drive
- Disk drives have peculiar characteristics
  - Long, and worse, variable access latencies
  - Accesses performed in chunks of fixed size
    - Atomicity only for accesses of that size
  - Highly variable performance depending on exactly what gets put where
  - Unpleasant failure modes
- So the operating system needs to smooth out these oddities



# What Does That Lead To?

- Great effort by file system component of OS to put things in the right place on a disk
- Reordering of disk operations to improve performance
  - Which complicates providing atomicity
- Optimizations based on caching and read-ahead
  - Which complicates maintaining consistency
- Sophisticated organizations to handle failures

# Abstractions of Interpreters

- An interpreter is something that performs commands
- Basically, the element of a computer (abstract or physical) that gets things done
- At the physical level, we have a processor
- That level is not easy to use
- The OS provides us with higher level interpreter abstractions

# Basic Interpreter Components

- An instruction reference
  - Tells the interpreter which instruction to do next
- A repertoire
  - The set of things the interpreter can do
- An environment reference
  - Describes the current state on which the next instruction should be performed
- Interrupts
  - Situations in which the instruction reference pointer is overridden

# For Example,

- A CPU
- It has a program counter register indicating where the next instruction can be found
  - An instruction reference
- It supports a set of instructions
  - Its repertoire
- It has contents in registers and RAM
  - Its environment

# Another Example

- A process
- The OS maintains a program counter for the process
  - An instruction reference
- Its source code specifies its repertoire
- Its stack, heap, and register contents are its environment
  - With the OS maintaining pointers to all of them
- No other interpreters should be able to mess up the process' resources

# Implementing the Process Abstraction in the OS

- Easy if there's only one process
- But there almost always are multiple processes
- The OS has a certain amount of physical memory
  - To hold the environment information
- There is usually only one set of registers
- The process doesn't have exclusive access to the CPU
  - Due to other processes

# What Does That Lead To?

- Schedulers to share the CPU among various processes
- Memory management hardware and software
  - To multiplex memory use among the processes
  - Giving each the illusion of full exclusive use of memory
- Access control mechanisms for other memory abstractions
  - So other processes can't fiddle with my files

# Abstractions of Communications Links

- A communication link allows one interpreter to talk to another
  - On the same or different machines
- At the physical level, wires and cables
- At more abstract levels, networks and interprocess communication mechanisms
- Some similarities to memory abstractions
  - But also differences



# Basic Communication Link Operations

- SEND(link\_name, outgoing\_message\_buffer)
  - Send some information contained in the buffer on the named link
- RECEIVE(link\_name, incoming\_message\_buffer)
  - Read some information off the named link and put it into the buffer
- Like WRITE and READ, in some respects

# Why Are Communication Links Distinct From Memory?

- Highly variable performance
- Potentially hostile environment for the operations
- Generally asynchronous
- Receiver may only perform the operation because the SEND occurred
  - Unlike a typical READ
- No necessary guarantee of delivery

# An Example Communications Link

- A Unix-style socket
- SEND interface:
  - `send(int sockfd, const void *buf, size_t len, int flags)`
  - The `sockfd` is the link name
  - The `buf` is the outgoing message buffer
- RECEIVE interface:
  - `recv(int sockfd, void *buf, size_t len, int flags)`
  - Same parameters as for `send`

# What About Those Other Socket Parameters?

- The `len` and `flag` fields?
- A common attribute of instances of abstractions
  - Especially higher level versions
- They provide additional semantics specific to the abstraction
- Generally improving the power of the higher level abstraction

# Implementing the Communications Link Abstraction in the OS

- A bit trickier than the memory and interpreter abstraction, in some cases
- Unlike those, the OS does not have full control of what's going on
- The network doesn't belong to the OS
  - Only its own network interface does
- Another entity is often doing half the work
  - Typically another machine's OS

# What Are the Implications?

- Greater uncertainty about the outcome of an operation
  - Things fail for reasons our OS can't see or learn
- Greater asynchrony
  - The remote OS might not regard the operations as equally important as our OS does
- Higher possibilities for security problems
  - Remote OS not equally trusted
  - Network between the two potentially untrustworthy

# What Do We Do About Those Issues?

- OS must be prepared for likely failures
- And high degrees of asynchrony
  - Bad idea to block entire system while waiting for the network
- OS shouldn't have complete trust in what comes in from the network
  - But often the OS is in no position to determine its trustworthiness

# System Services for OSes

- One major role of an operating system is providing services
  - To human users
  - To applications
- What services should an OS provide?



# An Object Oriented View of OS System Services

- Services are delivered through objects
  - Can be instantiated, named, and destroyed
  - They have specified properties
  - They support specified methods
- To understand a service, study its objects
  - How they are instantiated and managed
  - How client refers to them (names/handles)
  - What a client can do with them (methods)
  - How objects behave (interface specifications)

# Typical OS System Service Types

- Execution objects
  - Processes, threads, timers, signals
- Data objects
  - Files, devices, segments, file systems
- Communications objects
  - Sockets, messages, remote procedure calls
- Protection objects
  - Users, user groups, process groups
- Naming objects
  - Directories, DNS domains, registries

# System Services and Abstractions

- Services are commonly implemented by providing appropriate abstractions
- For example,
  - The service of allowing user code to run in a computing environment
  - Requires a couple of abstractions, at least:
    - The virtual environment abstraction
    - The process abstraction

# The Virtual Environment

## Abstraction

- A CPU executes one program at a time
  - It is a serially reusable resource
- But we want to run multiple programs “simultaneously”
  - Without them treading on each other’s toes
- A good way to do that is to build a virtual execution environment abstraction
  - Make it look like each program has its own computer

# What Should This Abstraction Provide?

- Each program should see its own resource set
  - A complete virtual computer with all elements
    - CPU
    - Memory
    - Persistent storage
    - Peripherals
- Isolation from other activities
  - Including non-related OS activities
- Each program should think it has the real machine to itself

# How To Do That?

- We won't go into detail now
  - But will later
- In essence, the OS must multiplex its real resources
  - Among the various process' virtual computers
- Requiring care in saving and restoring state
- And attention to fair use and processes' various performance requirements

# The Process Service

- Given we want per program virtual environments,
- We need an interpreter abstraction that provides the ability to run user code
  - The process
- With some very useful properties:
  - Isolation from other code
  - Isolation from many system failures
  - Guarantees of access to certain resources
- Processes can communicate and coordinate
  - But do so through the OS
  - Which provides isolation and synchronization

# What Is a Process?

- An interpreter that executes a single program
  - It provides illusion of continuous execution
  - Despite fact that the actual CPU is time-shared
    - Runs process A, then process B, then process A
- What virtual environment does a program see?
  - Programs don't run on a real bare computer
  - They run inside of a process
  - Process state is saved when it is not running
  - Process state is restored when it runs again



# Processes and Programs

- Program = set of executable instructions
  - Many processes can run the same program
- Process = executing instance of program
  - It has saved state
    - Memory, contents, program counter, registers, ...
  - It has resources and privileges
    - Open files, user-ID, capabilities, ...
  - It may be the unit of CPU sharing
    - CPU runs one process, then another

# Problems With the Process Abstraction

- Processes are very expensive
  - To create: they own resources
  - To dispatch: they have address spaces
- Different processes are very distinct
  - They cannot share the same address space
  - They cannot (usually) share resources
- Not all programs want strong separation
  - Cooperating parallel threads of execution
  - All are trusted because they run same code

# So the Process Abstraction Isn't Sufficient

- To meet common user needs
- What if I have a program that can do multiple things simultaneously?
- And requires regular, cheap communications between those different things?
- Processes are too expensive
- And make regular communications costly
- So I need another abstraction

# Threads

- An abstraction built on top of the process abstraction
- Each process contains one or more threads
- Each thread has some separate context of its own
  - Like a program counter and scheduling info
- But otherwise shares the resources of its process
- Threads within a process can thus communicate easily and cheaply

# Characteristics of Threads

- Strictly a unit of execution/scheduling
  - Each thread has its own stack, PC, registers
- Multiple threads can run in a process
  - They all share the same code and data space
  - They all have access to the same resources
  - This makes the cheaper to create and run
- Sharing the CPU between multiple threads
  - User level threads (with voluntary yielding)
  - Kernel threads (with preemption)

# Using the Abstractions

- When a programmer wants to run code, then, he can choose between abstractions
- Does he want just a process?
- Or does he want a process containing multiple threads?
- Or perhaps multiple processes?
  - With one thread each?
  - With multiple threads?

# When To Use Processes

- When running multiple distinct programs
- When creation/destruction are rare events
- When running programs (even instances of the same code) with distinct privileges
- When there are limited interactions and few shared resources
- When you need to prevent interference between programs
  - Or need to protect one from failures of the other

# An Example of Choosing Processes

- When implementing compilation in a shell script

***cpp \$1.c | cc1 | ccopt > \$1.s***

***as \$1.s***

***ld /lib/crt0.o \$1.o /lib/libc.so***

***mv a.out \$1***

***rm \$1.s \$1.o***

- Each of these programs gets a separate process



# Why?

- The activities are serial
- The only resources to be shared are through the file system
- Failure of one program could damage the others if too much is shared
  - Who knows what `rm` might get rid of, for example?

# When To Use Threads

- When there are parallel activities in a single program
- When there will be frequent creation and destruction
- When all activities can run with same privileges
- When they need to share resources
- When they exchange many messages/signals
- When there's no need to protect them from each other

# An Example for Choosing Threads

- A web server
- Multiple users will request service
- Desirable to share much of the server data
  - Such as copies of pages many users want to see
  - And information about overall load and performance
- But the pages can be served to users in parallel
  - In particular, if serving one user's page is slow, don't slow down other users

# Which Abstraction To Choose?

- If you use multiple processes
  - Your application may run much more slowly
  - It may be difficult to share some resources
- If you use multiple threads
  - You will have to create and manage them
  - You will have to serialize resource use
  - Your program will be more complex to write
  - You may get weird bugs
- TANSTAAFL

– There Ain't No Such Thing As A Free Lunch

# Generalizing the Concepts

- There are many other abstractions offered by the OS
- Often they provide different ways of achieving similar goals
  - Some higher level, some lower level
- The OS must do work to provide each abstraction
  - The higher level, the more work
- Programmers and users have to choose the right abstractions to work with

# Abstractions and Layering

- It's common to create increasingly complex services by layering abstractions
  - E.g., a file system layers on top of an abstract disk, which layers on top of a real disk
- Layering allows good modularity
  - Easy to build multiple services on a lower layer
    - E.g., multiple file systems on one disk
  - Easy to use multiple underlying services to support a higher layer
  - E.g., file system can have either a single disk or a RAID below it

# A Downside of Layering

- Layers typically add performance penalties
- Often expensive to go from one layer to the next
  - Since it frequently requires changing data structures or representations
  - At least involves extra instructions
- Another downside is that lower layer may limit what the upper layer can do
  - E.g., an abstract disk prevents disk operation reorderings to maximize performance

# Layer Bypassing

- Often necessary to allow a high layer to access much lower layers
  - Not going through one or more intermediaries
- Most commonly for performance reasons
- If the higher layer plans to use the very low level layer's services,
  - Why pay the cost of the intermediate layer?
- Has its downsides, too
  - Intermediate layer can't help or understand